

```
import torch
import numpy as np
import torch.nn as nn
import torch.optim as optim
from openpyxl.styles.builtins import output
from torch.nn import CrossEntropyLoss
from torch.utils.data import DataLoader, TensorDataset
from sklearn.metrics import accuracy_score, confusion_matrix
import matplotlib.pyplot as plt
from torchvision import datasets
from torchvision.transforms import ToTensor
import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
import torchvision
import torchvision.transforms as transforms
import torchvision.models as models
```

```
from google.colab import drive
drive.mount('/content/drive')
```

↗ Mounted at /content/drive

```
!pip install datasets
```

↗ Collecting datasets

```
  Downloading datasets-3.2.0-py3-none-any.whl.metadata (20 kB)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from datasets) (3.16.1)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from datasets) (1.26.4)
Requirement already satisfied: pyarrow>=15.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (17.0.0)
Collecting dill<0.3.9,>=0.3.0 (from datasets)
  Downloading dill-0.3.8-py3-none-any.whl.metadata (10 kB)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from datasets) (2.2.2)
```

Requirement already satisfied: requests>=2.32.2 in /usr/local/lib/python3.10/dist-packages (from datasets) (2.32.2)
Requirement already satisfied: tqdm>=4.66.3 in /usr/local/lib/python3.10/dist-packages (from datasets) (4.67.1)
Collecting xxhash (from datasets)

Downloading xxhash-3.5.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (12 kB)

Collecting multiprocess<0.70.17 (from datasets)

Downloading multiprocess-0.70.16-py310-none-any.whl.metadata (7.2 kB)

Collecting fsspec<=2024.9.0,>=2023.1.0 (from fsspec[http]<=2024.9.0,>=2023.1.0->datasets)

Downloading fsspec-2024.9.0-py3-none-any.whl.metadata (11 kB)

Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets) (3.11.11)

Requirement already satisfied: huggingface-hub>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.23.0)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from datasets) (24.2)

Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from datasets) (6.0.2)

Requirement already satisfied: aiohappyeyeballs>=2.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (2.4.4)

Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.3.1)

Requirement already satisfied: async-timeout<6.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (4.0.3)

Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (23.2.0)

Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.4.1)

Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (6.0.5)

Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (0.2.0)

Requirement already satisfied: yarl<2.0,>=1.17.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.17.0)

Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (4.11.0)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.32.2->datasets) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.32.2->datasets) (3.10.1)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.32.2->datasets) (2.2.3)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.32.2->datasets) (2024.7.4)

Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2.9.0.post0)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2024.2)

Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2024.2)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->datasets) (1.17.0)

Downloading datasets-3.2.0-py3-none-any.whl (480 kB)

480.6/480.6 kB 12.8 MB/s eta 0:00:00

Downloading dill-0.3.8-py3-none-any.whl (116 kB)

116.3/116.3 kB 8.7 MB/s eta 0:00:00

Downloading fsspec-2024.9.0-py3-none-any.whl (179 kB)

179.3/179.3 kB 12.6 MB/s eta 0:00:00

Downloading multiprocess-0.70.16-py310-none-any.whl (134 kB)

134.8/134.8 kB 12.6 MB/s eta 0:00:00

Downloading xxhash-3.5.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (194 kB)

194.1/194.1 kB 17.4 MB/s eta 0:00:00

Installing collected packages: xxhash, fsspec, dill, multiprocessing, datasets

Attempting uninstall: fsspec

Found existing installation: fsspec 2024.10.0

Uninstalling fsspec-2024.10.0:

Successfully uninstalled fsspec-2024.10.0

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This gcfs 2024.10.0 requires fsspec==2024.10.0, but you have fsspec 2024.9.0 which is incompatible.

Successfully installed datasets-3.2.0 dill-0.3.8 fsspec-2024.9.0 multiprocessing-0.70.16 xxhash-3.5.0

```
df = pd.read_csv('/content/drive/MyDrive/Mental project/dataset/Combined Data.csv')
```

```
print(df.shape)
```

```
df.head()
```

↔ (53043, 3)

Unnamed: 0		statement	status	
0	0	oh my gosh	Anxiety	
1	1	trouble sleeping, confused mind, restless hear...	Anxiety	
2	2	All wrong, back off dear, forward doubt. Stay ...	Anxiety	
3	3	I've shifted my focus to something else but I'...	Anxiety	
4	4	I'm restless and restless, it's been a month n...	Anxiety	

Next steps:

[Generate code with df](#)

[View recommended plots](#)

[New interactive sheet](#)

```
df = df[['statement', 'status']]
```

```
df.head()
```



	statement	status
--	-----------	--------



0	oh my gosh	Anxiety
1	trouble sleeping, confused mind, restless hear...	Anxiety
2	All wrong, back off dear, forward doubt. Stay ...	Anxiety
3	I've shifted my focus to something else but I'...	Anxiety
4	I'm restless and restless, it's been a month n...	Anxiety



Next steps:

[Generate code with df](#)[View recommended plots](#)[New interactive sheet](#)

```
df.isnull().sum()
```



	0
--	---

statement	362
-----------	-----

status	0
--------	---

dtype: int64

```
df.dropna(inplace=True)
```

```
df.isnull().sum()
```

↔ 0

statement	0
status	0

dtype: int64

```
df.duplicated().sum()
```

↔ 1588

```
df.drop_duplicates(inplace=True, keep='first')
```

```
df.duplicated().sum()
```

↔ 0

```
df2 = df.sample(6000).reset_index(drop=True)
df2.shape
```

↔ (6000, 2)

```
df2['status'].value_counts()
```



count	
status	
Normal	1827
Depression	1777
Suicidal	1269
Anxiety	428
Bipolar	309
Stress	283
Personality disorder	107

dtype: int64

```
from imblearn.over_sampling import RandomOverSampler
```

```
ros = RandomOverSampler(random_state=42)
```

```
X_resampled, y_resampled = ros.fit_resample(df2[['statement']], df2['status'])
```

```
df_resampled = pd.concat([X_resampled, y_resampled], axis=1)
```

```
df_resampled.head()
```



statement **status**



0	I will just introduce a bit.I am 21 years old....	Suicidal
1	Ikea king â	Normal
2	Health Anxiety, what has worked for me - Hopef...	Anxiety
3	blame society Blame moviesBlames stereotypesBl...	Depression
4	Are you a liberal activist?	Normal



Next steps:

[Generate code with df_resambled](#)

☒ [View recommended plots](#)

[New interactive sheet](#)

```
df_resambled['status'].value_counts()
```



count

status

Suicidal	1827
Normal	1827
Anxiety	1827
Depression	1827
Bipolar	1827
Stress	1827
Personality disorder	1827

dtype: int64

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()  
df_resambled['status'] = le.fit_transform(df_resambled['status'])  
df_resambled.head()
```

```
df_resambled['status'].value_counts()
```



count

status

6 1827

3 1827

0 1827

2 1827

1 1827

5 1827

4 1827

dtype: int64

```
from nltk.corpus import stopwords  
import nltk  
from nltk.tokenize import word_tokenize  
import string
```

```
nltk.download('punkt_tab')  
nltk.download('stopwords')
```

```
def cleaned_text(text):  
    text = text.lower()  
    token = word_tokenize(text)
```



```
stop_words = set(stopwords.words('english'))
filtered_sentence = [word for word in token if word not in stop_words and string.punctuation and word.isalnum()]
return " ".join(filtered_sentence)
```

```
text = "I am learning NLP ### AA @@@ !!! , any one can HEPL ME OUT ??? "
cleaned_text(text)
```

```
🔄 [nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt_tab.zip.
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
'learning nlp aa one hepl'
```

```
X = df_resambled['statement'].apply(cleaned_text)
y = df_resambled['status']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
from transformers import BertTokenizer, BertForSequenceClassification, AutoTokenizer, AutoModelForSequenceClassification
```

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
train_tokenizer = tokenizer(list(X_train), padding=True, truncation=True, max_length=128)
test_tokenizer = tokenizer(list(X_test), padding=True, truncation=True, max_length=128)
```

```

➦ /usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens)
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(


tokenizer_config.json: 100% 48.0/48.0 [00:00<00:00, 1.19kB/s]

vocab.txt: 100% 232k/232k [00:00<00:00, 2.90MB/s]

tokenizer.json: 100% 466k/466k [00:00<00:00, 5.74MB/s]

config.json: 100% 570/570 [00:00<00:00, 16.8kB/s]

```



```
from datasets import Dataset
```

```
train_dataset = Dataset.from_dict({'input_ids': train_tokenizer['input_ids'], 'attention_mask': train_tokenizer['attention_mask']})
test_dataset = Dataset.from_dict({'input_ids': test_tokenizer['input_ids'], 'attention_mask': test_tokenizer['attention_mask']})
```

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
device
```

```
➦ device(type='cuda')
```

```
num_labels = len(df_resampled['status'].unique())
model = BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels=num_labels)
```

```
training_args = TrainingArguments(
    output_dir="./results",          # Output directory for results
    evaluation_strategy="epoch",      # Evaluate once per epoch
    save_strategy="epoch",           # Save model at the end of each epoch to match evaluation strategy
    learning_rate=2e-5,              # Learning rate
    per_device_train_batch_size=16,   # Batch size for training
    per_device_eval_batch_size=16,    # Batch size for evaluation
)
```

```

num_train_epochs=5,          # Increase number of epochs
weight_decay=0.01,          # Strength of weight decay
logging_dir="./logs",        # Directory for logging
logging_steps=10,           # Log every 10 steps
lr_scheduler_type="linear",  # Use linear learning rate scheduler with warmup
warmup_steps=500,           # Number of warmup steps for learning rate scheduler
load_best_model_at_end=True, # Load the best model at the end of training
metric_for_best_model="eval_loss", # Monitor eval loss to determine the best model
save_total_limit=3,          # Limit the number of checkpoints to save
gradient_accumulation_steps= 2 # Simulate larger batch size if GPU memory is limite
)

trainer = Trainer(
    model = model.to(device),
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=test_dataset,
)

trainer.train()

```



model.safetensors: 100%

440M/440M [00:07<00:00, 72.2MB/s]

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
/usr/local/lib/python3.10/dist-packages/transformers/training_args.py:1575: FutureWarning: `evaluation_strategy` is deprecated in favor of `eval_strategy`
warnings.warn()

wandb: **WARNING** The `run_name` is currently set to the same value as `TrainingArguments.output_dir`. If this was not intended, please set `run_name` to a unique name.

wandb: Using wandb-core as the SDK backend. Please refer to <https://wandb.me/wandb-core> for more information.

wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: <https://wandb.me/wandb-server>)

wandb: You can find your API key in your browser here: <https://wandb.ai/authorize>

wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit:

wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc

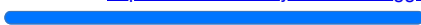
Tracking run with wandb version 0.19.1

Run data is saved locally in /content/wandb/run-20250112_111429-0h1q2aov

Syncing run [./results](#) to [Weights & Biases \(docs\)](#)

View project at <https://wandb.ai/rickydoan144/huggingface>

View run at <https://wandb.ai/rickydoan144/huggingface/runs/0h1q2aov>

 [1600/1600 22:09, Epoch 5/5]

Epoch	Training Loss	Validation Loss
-------	---------------	-----------------

1	2.058000	0.965442
---	----------	----------

2	1.025200	0.476821
---	----------	----------

3	0.407700	0.316571
---	----------	----------

4	0.346100	0.285863
---	----------	----------

5	0.239400	0.295584
---	----------	----------

```
TrainOutput(global_step=1600, training_loss=1.1807581675052643, metrics={'train_runtime': 1361.134, 'train_samples_per_second': 37.583, 'train_steps_per_second': 1.175, 'total_flos': 3365012567904000.0, 'train_loss': 1.1807581675052643, 'epoch': 5.0})
```

```
from sklearn.metrics import classification_report, confusion_matrix
```

```
predictions, labels, _ = trainer.predict(test_dataset)
```

```

predictions_label = np.argmax(predictions, axis=1)
print(classification_report(labels, predictions_label, target_names= le.classes_))

```




	precision	recall	f1-score	support
Anxiety	0.95	0.99	0.97	357
Bipolar	0.95	0.99	0.97	377
Depression	0.77	0.65	0.71	373
Normal	0.94	0.89	0.92	374
Personality disorder	0.99	1.00	1.00	346
Stress	0.98	0.99	0.98	330
Suicidal	0.74	0.82	0.78	401
accuracy			0.90	2558
macro avg	0.91	0.91	0.90	2558
weighted avg	0.90	0.90	0.90	2558

```

cm = confusion_matrix(labels, predictions_label)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=le.classes_, yticklabels=le.classes_)
plt.xlabel('Predicted')
plt.ylabel('True')

```

Text(50.7222222222221, 0.5, 'True')



Generated code may be subject to a license | [stevenwchien/csml-iw-rrp](#)

```
trainer.save_model('/content/drive/MyDrive/Mental project/model_bert_mental')
```

```
tokenizer.save_pretrained('/content/drive/MyDrive/Mental project/model_bert_mental')
```

```
model = AutoModelForSequenceClassification.from_pretrained('/content/drive/MyDrive/Mental project/model_bert_mental2')
# tokenizer = AutoTokenizer.from_pretrained('/content/drive/MyDrive/Mental project/model_bert_mental')
tokenizer = AutoTokenizer.from_pretrained('bert-base-uncased')
```

```
from joblib import dump
```

```
dump(le, '/content/drive/MyDrive/Mental project/label_encoder.joblib')
```

```
↔ ['/content/drive/MyDrive/Mental project/label_encoder.joblib']
```

```
model
```

```
↔ Show hidden output
```

```
def predict_sentiment(text):
    text_cleaned = cleaned_text(text)
    inputs = tokenizer(text_cleaned, padding=True, truncation=True, max_length=128, return_tensors="pt")
    # Move input tensors to the same device as the model
    inputs = {key: val.to(device) for key, val in inputs.items()}
    outputs = model(**inputs)
```